

NAAMIIKA

VOICE ASSISTANT SYSTEM

Complete Technical Documentation

Document Version: v2.1.7 (v2_demo_22Jan)

Classification: Internal Technical Reference

Generated: January 27, 2026

Repository: github.com/bhaveshbakshi633/indu_v2_1_2

Branch: v2_demo_22Jan

Prepared by: Engineering Team

Organization: SS Innovations International

TABLE OF CONTENTS

Part I: System Overview

1. Executive Summary
2. System Architecture

Part II: Voice Processing Pipeline

1. Voice Activity Detection (VAD)
2. Speech-to-Text (STT)
3. Intent Detection & Reasoning
4. Action Registry & Safety

Part III: AI & Response Generation

1. RAG-Enhanced LLM Processing
2. Text-to-Speech (TTS)

Part IV: Robot Control

1. SLAM Navigation System
2. G1 Robot Integration

Part V: Operations

1. Configuration System
2. API Reference
3. Deployment Guide
4. Troubleshooting

Appendix A. Complete File Reference B. Configuration Values

CHAPTER 1: EXECUTIVE SUMMARY

1.1 Document Purpose

This document provides comprehensive technical documentation for the NAAMIKA Voice Assistant System, a voice-controlled interface for the Unitree G1 humanoid robot. It serves as the authoritative reference for system architecture, code implementation, deployment procedures, and operational guidelines.

1.2 System Overview

NAAMIKA is a distributed voice-controlled humanoid robot assistant designed for the Unitree G1 platform. The system integrates multiple AI components to enable natural voice interaction and physical robot control.

Core Capabilities:

Capability	Technology	Status
Voice Recognition	Whisper STT (GPU)	Operational
Natural Language Understanding	Ollama LLM + RAG	Operational
Text-to-Speech	Chatterbox TTS	Operational
Intent Classification	Pattern Matching + LLM	Operational
Emergency Stop	Fast-path Detection	Operational
Gesture Control	HTTP Action Bridge	Operational
Locomotion	Time-based Motion	Operational
SLAM Navigation	Unitree SDK	Operational
Face Recognition	Shakal	Operational

1.3 Architecture Summary

The system operates across four networked machines:

Machine	Role	IP Address
This PC (ssi)	Brain Server, TTS	172.16.6.19
G1 Robot	Motion Control, Audio	172.16.2.242
Isaac PC	LLM Inference	172.16.4.226
New PC (b)	STT Inference	172.16.4.250

1.4 Critical Design Principles

Safety-First Architecture:

- All physical actions require FSM state validation
- MEDIUM/HIGH risk actions require explicit user confirmation
- Emergency STOP bypasses all processing pipelines instantly
- 22-action whitelist prevents unauthorized command execution

Distributed Processing:

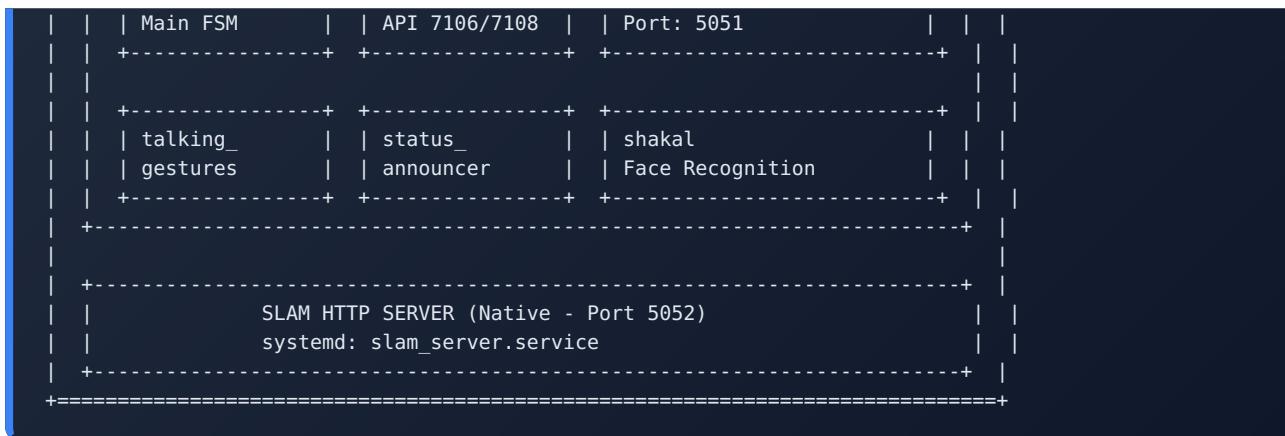
- Compute-intensive tasks (LLM, STT) offloaded to GPU servers
- HTTP-based communication for reliability over ROS2
- Graceful degradation with fallback mechanisms

CHAPTER 2: SYSTEM ARCHITECTURE

2.1 High-Level System Diagram

The following diagram illustrates the complete system architecture including all four machines and their interconnections.





2.2 Network Configuration

Service	Machine	IP Address	Port	Protocol
Voice Assistant (Brain)	This PC	172.16.6.19	8080	HTTPS
Chatterbox TTS	This PC	172.16.6.19	8000	HTTP
Whisper STT	New PC (b)	172.16.4.250	8001	HTTP
Ollama LLM	Isaac PC	172.16.4.226	11434	HTTP
G1 Audio Receiver	G1 Robot	172.16.2.242	5050	HTTP
G1 HTTP Action Bridge	G1 Robot	172.16.2.242	5051	HTTP
G1 SLAM Server	G1 Robot	172.16.2.242	5052	HTTP

2.3 Data Flow Pipeline

The voice processing pipeline consists of six stages from user speech to robot action or response.



CHAPTER 3: VOICE ACTIVITY DETECTION PIPELINE

3.1 Overview

The VAD pipeline detects when a user starts and stops speaking using the Silero VAD neural network model. This enables hands-free voice interaction without push-to-talk buttons.

File Location: `vad_pipeline/`

3.2 VAD Processor

File: `vad_pipeline/vad_processor.py`

3.2.1 Configuration Parameters

```
class VADProcessor:  
    def __init__(  
        self,  
        threshold: float = 0.5,           # Speech confidence threshold  
        frame_count: int = 3,            # Frames for confirmation  
        sample_rate: int = 16000,          # Audio sample rate (Hz)  
        chunk_size: int = 512,             # Samples per chunk  
        interrupt_threshold: float = 1.0,   # Strict interrupt threshold  
        interrupt_frames: int = 5         # Frames for interrupt  
    ):
```

3.2.2 Core Methods

Single Frame Detection:

```
def is_speech(self, audio_chunk: np.ndarray) -> Tuple[bool, float]:  
    """  
    Process single audio chunk through Silero VAD.  
    Returns: (is_speech: bool, probability: float)  
    """  
    audio_tensor = torch.from_numpy(audio_chunk).float()  
    speech_prob = self.model(audio_tensor, self.sample_rate).item()  
    return speech_prob >= self.threshold, speech_prob
```

Multi-Frame Speech Start Detection:

```
def detect_speech_start(self, audio_chunk) -> bool:
    """
    Requires multiple consecutive frames above threshold.
    Prevents false positives from transient noise.
    """
    is_speech, prob = self.is_speech(audio_chunk)
    self.frame_buffer.append(prob)

    if len(self.frame_buffer) >= self.frame_count:
        recent = self.frame_buffer[-self.frame_count:]
        return all(p >= self.threshold for p in recent)
    return False
```

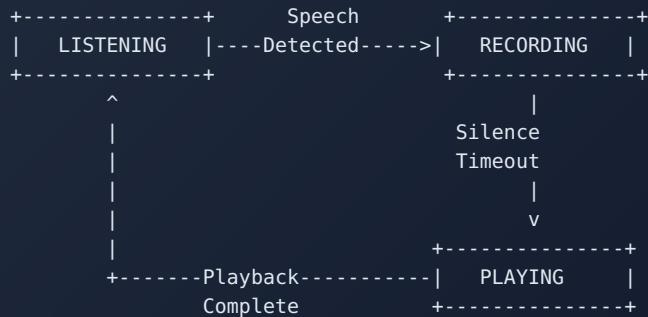
Interrupt Detection During Playback:

```
def detect_interrupt(self, audio_chunk) -> bool:
    """
    Used during TTS playback to detect user interruption.
    Uses stricter threshold across multiple frames.
    """
    is_speech, prob = self.is_speech(audio_chunk)
    self.interrupt_buffer.append(prob)

    if len(self.interrupt_buffer) >= self.interrupt_frames:
        high_conf = sum(1 for p in self.interrupt_buffer
                        if p >= self.interrupt_threshold)
        return high_conf >= (0.67 * self.interrupt_frames)
    return False
```

3.3 Pipeline State Machine

The pipeline manager operates as a three-state machine:



3.4 Audio Recorder

File: `vad_pipeline/audio_recorder.py`

```
class AudioRecorder:  
    def __init__(  
        self,  
        vad_processor: VADProcessor,  
        sample_rate: int = 16000,  
        silence_timeout: float = 1.5,    # Stop after 1.5s silence  
        max_duration: float = 30.0      # Maximum recording length  
    ):  
  
        def record_from_monitor(self, audio_monitor) -> bytes:  
            """Record audio until silence detected."""  
            buffer = []  
            silence_start = None  
  
            while True:  
                chunk = audio_monitor.get_chunk()  
                is_speech, prob = self.vad.is_speech(chunk)  
                buffer.append(chunk)  
  
                if is_speech:  
                    silence_start = None  
                else:  
                    if silence_start is None:  
                        silence_start = time.time()  
                    elif time.time() - silence_start > self.silence_timeout:  
                        break  
  
                if self.get_buffer_duration() > self.max_duration:  
                    break  
  
            return np.concatenate(buffer).tobytes()
```

CHAPTER 4: SPEECH-TO-TEXT PROCESSING

4.1 Whisper Server Architecture

File: `whisper_server.py` Deployment: 172.16.4.250:8001

4.1.1 Server Configuration

```
MODELS_AVAILABLE = {  
    "tiny": "tiny",      # 39M params, 1GB VRAM  
    "base": "base",     # 74M params, 2GB VRAM  
    "small": "small",   # 244M params, 3GB VRAM  
    "medium": "medium", # 769M params, 5GB VRAM (DEFAULT)  
    "large": "large-v3" # 1.5B params, 10GB VRAM  
}  
  
DEFAULT_MODEL = "medium"  
DEVICE = "cuda"  
COMPUTE_TYPE = "float16"  
PORT = 8001
```

4.1.2 Transcription Endpoint

```
@app.post("/transcribe")
async def transcribe(
    file: UploadFile,
    model: str = "medium",
    language: Optional[str] = None
):
    """
    Process audio file and return transcript.

    Request:
    - file: WAV audio file
    - model: tiny/base/small/medium/large
    - language: en, hi, es, etc. (auto-detect if None)

    Response:
    {
        "text": "transcribed text",
        "language": "en",
        "model": "medium",
        "duration_seconds": 5.2,
        "processing_time_seconds": 0.8
    }
    """
    whisper_model = load_model(model)

    segments, info = whisper_model.transcribe(
        audio_path,
        language=language,
        task="transcribe",
        beam_size=5
    )

    text = " ".join([s.text for s in segments]).strip()

    return TranscribeResponse(
        text=text,
        language=info.language,
        model=model,
        duration_seconds=info.duration,
        processing_time_seconds=time.time() - start
    )
```

4.2 Brain Server STT Integration

File: `server.py`

```
async def transcribe_audio(audio_bytes: bytes) -> str:  
    """Send audio to Whisper server and get transcript."""  
    config = load_config()  
    stt_config = config["stt"]["whisper_server"]  
  
    url = f"http://{{stt_config['host']}:{stt_config['port']}}/transcribe"  
  
    async with httpx.AsyncClient(timeout=30.0) as client:  
        files = {"file": ("audio.wav", audio_bytes, "audio/wav")}  
        params = {  
            "model": stt_config.get("model", "medium"),  
            "language": stt_config.get("language", "en")  
        }  
  
        response = await client.post(url, files=files, params=params)  
        response.raise_for_status()  
        return response.json()["text"]
```

CHAPTER 5: INTENT DETECTION & REASONING

5.1 Overview

The Intent Reasoner determines whether user speech is an action command or conversational input. It uses a two-tier approach: fast pattern matching first, LLM classification as fallback.

File: `intent_reasoner.py`

5.2 Intent Types and Results

```
class IntentType(Enum):
    ACTION = "action"          # Physical robot action
    CONVERSATION = "conversation"  # LLM response needed
    QUERY = "query"            # Information request

@dataclass
class IntentResult:
    intent_type: IntentType
    action_name: Optional[str]   # e.g., "WAVE", "FORWARD"
    confidence: float           # 0.0 to 1.0
    requires_confirmation: bool  # MEDIUM/HIGH risk
    reason: str                 # Human-readable explanation
    original_transcript: str    # Original STT output
```

5.3 Action Pattern Dictionary

```
ACTION_PATTERNS = {  
    # System actions (HIGH priority)  
    "initialize": ("INIT", 0.95),  
    "boot up": ("INIT", 0.95),  
    "chalu karo": ("INIT", 0.95),      # Hindi  
  
    "ready": ("READY", 0.90),  
    "get ready": ("READY", 0.95),  
    "taiyaar": ("READY", 0.90),       # Hindi  
  
    "damp": ("DAMP", 0.95),  
    "relax": ("DAMP", 0.85),  
  
    # Posture actions  
    "stand up": ("STANDUP", 0.95),  
    "utho": ("STANDUP", 0.95),        # Hindi  
    "sit down": ("SIT", 0.95),  
  
    # Motion actions  
    "walk forward": ("FORWARD", 0.95),  
    "aage": ("FORWARD", 0.90),        # Hindi  
    "walk back": ("BACKWARD", 0.95),  
    "turn left": ("LEFT", 0.95),  
    "turn right": ("RIGHT", 0.95),  
  
    # Gestures  
    "wave": ("WAVE", 0.95),  
    "shake hand": ("SHAKE_HAND", 0.95),  
    "hug": ("HUG", 0.95),  
    "namaste": ("namaste1", 0.95),  
  
    # SLAM  
    "start mapping": ("START_MAPPING", 0.95),  
    "stop mapping": ("STOP_MAPPING", 0.95),  
    "start navigation": ("START_NAVIGATION", 0.95),  
    "list waypoints": ("LIST_WAYPOINTS", 0.95),  
}
```

5.4 Local Pattern Matching

```
def parse_intent_local(transcript: str) -> IntentResult:
    """
    Fast local intent parsing without LLM.

    Priority:
    1. Check if question (always conversation)
    2. Check SLAM parameterized commands
    3. Check direct action patterns
    4. Default to conversation
    """
    lower = transcript.lower().strip()

    # Questions are always conversation
    if is_question(transcript):
        return IntentResult(
            intent_type=IntentType.CONVERSATION,
            action_name=None,
            confidence=0.95,
            requires_confirmation=False,
            reason="Question detected"
        )

    # Check SLAM parameterized commands
    waypoint = is_save_waypoint_command(transcript)
    if waypoint:
        return IntentResult(
            intent_type=IntentType.ACTION,
            action_name=f"SAVE WAYPOINT:{waypoint}",
            confidence=0.90,
            requires_confirmation=True,
            reason=f"Save waypoint: '{waypoint}'"
        )

    # Check direct pattern match
    match = match_action_pattern(transcript)
    if match:
        action_name, confidence = match
        needs_confirm = action_name in [
            "INIT", "DAMP", "ZERO_TORQUE",
            "STANDUP", "SIT", "FORWARD", "BACKWARD"
        ]
        return IntentResult(
            intent_type=IntentType.ACTION,
            action_name=action_name,
            confidence=confidence,
            requires_confirmation=needs_confirm,
            reason=f"Matched pattern '{action_name}'"
        )

    # Default to conversation
    return IntentResult(
        intent_type=IntentType.CONVERSATION,
        action_name=None,
        confidence=0.8,
        reason="No action pattern matched"
    )
```

5.5 Confirmation Flow

Confirmation Patterns

```
POSITIVE_PATTERNS = [
    "yes", "yeah", "yep", "ok", "okay", "sure", "confirm",
    "go ahead", "absolutely",
    "haan", "ha", "theek", "bilkul", "zaroor" # Hindi
]

NEGATIVE_PATTERNS = [
    "no", "nope", "cancel", "stop", "don't", "abort",
    "nahi", "mat", "ruk", "rehne do" # Hindi
]
```

IntentReasoner State Machine

```
class IntentReasoner:
    def __init__(self):
        self.pending_action = None
        self.awaiting_confirmation = False

    def process(self, transcript: str) -> IntentResult:
        # Check if awaiting confirmation first
        if self.awaiting_confirmation and self.pending_action:
            if is_confirmation(transcript):
                action = self.pending_action
                self.clear_pending()
                return IntentResult(
                    intent_type=IntentType.ACTION,
                    action_name=action,
                    confidence=1.0,
                    requires_confirmation=False,
                    reason="User confirmed"
                )
            elif is_rejection(transcript):
                self.clear_pending()
                return IntentResult(
                    intent_type=IntentType.CONVERSATION,
                    reason="User cancelled"
                )

        # Normal intent parsing
        result = parse_intent_local(transcript)

        # Store for confirmation if needed
        if result.requires_confirmation:
            self.pending_action = result.action_name
            self.awaiting_confirmation = True

    return result
```

CHAPTER 6: ACTION REGISTRY & SAFETY SYSTEM

6.1 Overview

The Action Registry defines all valid robot actions with their risk levels, FSM requirements, and API mappings. This serves as the single source of truth for action validation.

File: `action_registry.py`

6.2 Risk Classification

```
class RiskLevel(Enum):
    LOW = "low"      # Gestures - immediate execution
    MEDIUM = "medium" # Locomotion - needs confirmation
    HIGH = "high"     # System commands - extra validation

class ActionType(Enum):
    GESTURE = "gesture"
    MOTION = "motion"
    POSTURE = "posture"
    SYSTEM = "system"
    MODE = "mode"
    QUERY = "query"
```

6.3 FSM State Constants

```
# Unitree G1 FSM States
FSM_ZERO_TORQUE = 0    # Motors off - robot may fall
FSM_DAMP = 1            # Damping mode - gentle stop
FSM_SQUAT = 2           # Squat position
FSM_SIT = 3             # Sitting position
FSM_STANDUP = 4         # Standing from sit/lie
FSM_START = 500          # Sport/active mode
FSM_READY = 801          # Ready for arm control

# Valid FSM sets
FSM_ANY = {0, 1, 2, 3, 4, 500, 801}
FSM_STANDING = {500, 801}
FSM_READY_ONLY = {801}
```

6.4 Action Definitions (Selected)

Action	Type	Risk	Required FSM	API
WAVE	Gesture	LOW	801	7106:0
SHAKE_HAND	Gesture	LOW	801	7106:2
HUG	Gesture	LOW	801	7108:hug
NAMASTE	Gesture	LOW	801	7108:namaste1
FORWARD	Motion	MEDIUM	500, 801	7105
BACKWARD	Motion	MEDIUM	500, 801	7105
LEFT	Motion	MEDIUM	500, 801	7105
RIGHT	Motion	MEDIUM	500, 801	7105
STOP	Motion	LOW	Any	7105:stop
STANDUP	Posture	MEDIUM	1, 4	7101:4
SIT	Posture	MEDIUM	801	7101:3
INIT	System	HIGH	Any	orchestrator:init
DAMP	System	HIGH	Any	7101:1

6.5 Emergency Stop Fast Path

File: `stop_fast_path.py`

Stop Keywords

```
class StopType(Enum):
    EMERGENCY = "emergency" # "emergency", "emergency stop"
    IMMEDIATE = "immediate" # "stop", "halt", "freeze"
    CASUAL = "casual" # "ruk", "bas" (Hindi)

    EMERGENCY_KEYWORDS = ["emergency", "emergency stop"]
    IMMEDIATE_KEYWORDS = ["stop", "halt", "freeze", "stop now"]
    CASUAL_KEYWORDS = ["ruk", "ruk jao", "bas", "band karo"]

    # Exclusions (valid commands, not emergencies)
    STOP_EXCLUSIONS = ["stop mapping", "stop navigation", "stop talking"]
```

Detection Logic

```
def detect_stop(transcript: str) -> StopDetectionResult:
    """
    Check for STOP keywords BEFORE any other processing.
    Runs immediately after STT, bypasses LLM entirely.
    """
    lower = transcript.lower().strip()

    # First check exclusions
    for exclusion in STOP_EXCLUSIONS:
        if exclusion in lower:
            return StopDetectionResult(is_stop=False)

    # Priority: emergency > immediate > casual
    for stop_type, keywords in ALL_STOP_KEYWORDS.items():
        for keyword in keywords:
            if keyword in lower:
                return StopDetectionResult(
                    is_stop=True,
                    stop_type=stop_type,
                    matched_keyword=keyword
                )

    return StopDetectionResult(is_stop=False)
```

CHAPTER 7: RAG-ENHANCED LLM PROCESSING

7.1 Overview

The NAAMIIKA RAG Agent combines conversation memory with knowledge base retrieval for accurate, context-aware responses about SSi medical robotics.

File: `naamika_rag.py`

7.2 Architecture

```
User Query
  |
  v
+-----+
| STEP 1: Knowledge Retrieval (RAG)      |
|                                         |
| - Embed query with MiniLM-L6-v2       |
| - Search FAISS vector store           |
| - Retrieve top 5 relevant chunks     |
+-----+
  |
  v
+-----+
| STEP 2: Prompt Construction          |
|                                         |
| - CRITICAL_FACTS (anti-hallucination) |
| - Conversation history                |
| - RAG context                         |
| - User query                          |
+-----+
  |
  v
+-----+
| STEP 3: LLM Generation               |
|                                         |
| - Ollama with naamika:v1 model      |
| - Temperature: 0.2                   |
| - Context: 4096 tokens               |
+-----+
  |
  v
+-----+
| STEP 4: Response Processing         |
|                                         |
| - Save to conversation history      |
| - Return for TTS                    |
+-----+
```

7.3 Anti-Hallucination Facts

```
CRITICAL_FACTS = """
CRITICAL RULES - FOLLOW EXACTLY:
1. If you don't find information, say "I don't have that" - NEVER fabricate.
2. For competitors (da Vinci, Medtronic), redirect to SSi advantages.
3. You are NAAMIKA (NOT INDU), a humanoid robot. NOT the CEO.
4. NO MARKDOWN: Natural sentences only.
```

DR. SRIVASTAVA FACTS (USE EXACTLY):

- Education: J.L.N. Medical College, Ajmer
- Returned to India: 2011
- Awards: Golden Robot Surgical Award 2025
- World Record: 1,300+ beating heart TECAB surgeries

SSi STATS (December 2025):

- Installations: 168 worldwide
- Surgeries: 7,800+ performed
- Headquarters: Gurugram, India

SSi MANTRA SPECS:

- Components: FOUR (Surgeon Console, Arms, Vision, MUDRA)
- Degrees of Freedom: 7 DOF
- Scaling: 2:1, 3:1, 4:1
- Cost: Less than one-third of competitors

7.4 NaamikaAgent Class

```
class NaamikaAgent:  
    def __init__(  
        self,  
        knowledge_base_path: str = "naamika_knowledge_base.txt",  
        system_prompt_path: str = "naamika_system_prompt.txt",  
        model_name: str = "naamika:v1",  
        vector_store_path: str = "naamika_vectorstore",  
        ollama_host: str = "127.0.0.1",  
        ollama_port: int = 11434,  
        temperature: float = 0.2,  
        max_history: int = 10  
    ):  
        self.conversation_history = []  
  
        # Initialize embeddings  
        self.embeddings = HuggingFaceEmbeddings(  
            model_name="sentence-transformers/all-MiniLM-L6-v2"  
        )  
  
        # Load/create FAISS vector store  
        self.vectorstore = FAISS.load_local(vector_store_path, ...)  
  
        # Initialize retriever  
        self.retriever = self.vectorstore.as_retriever(  
            search_kwargs={"k": 5, "fetch_k": 20}  
        )  
  
        # Initialize LLM  
        self.llm = OllamaLLM(  
            model=model_name,  
            base_url=f"http://{{ollama_host}}:{ollama_port}",  
            temperature=temperature  
        )
```

7.5 Chat Method

```
def chat(self, user_query: str) -> str:  
    """Send query and get response with memory + RAG."""  
  
    # 1. RAG retrieval  
    rag_context = self._retrieve_context(user_query)  
  
    # 2. Format conversation history  
    chat_history = self._format_chat_history()  
  
    # 3. Build prompt with CRITICAL_FACTS  
    prompt = f"""{CRITICAL_FACTS}  
  
CONVERSATION HISTORY:  
{chat_history}  
  
RELEVANT KNOWLEDGE:  
{rag_context}  
  
USER: {user_query}  
  
RESPONSE:  
  
    # 4. Call LLM  
    response = self.llm.invoke(prompt).strip()  
  
    # 5. Save to history  
    self.conversation_history.append((user_query, response))  
  
    return response
```

CHAPTER 8: TEXT-TO-SPEECH PIPELINE

8.1 Overview

The TTS pipeline converts LLM responses to audio and streams to G1's speaker for natural voice output.

Files:

- `chatterbox_tts_client.py` - Client library
- External: Chatterbox server at 172.16.6.19:8000

8.2 Chatterbox TTS Client

```
class ChatterboxTTS:  
    def __init__(self, base_url: str = "http://172.16.6.19:8000"):  
        self.base_url = base_url  
  
    async def generate(  
        self,  
        text: str,  
        voice: str = None,  
        output_file: str = "output.wav"  
    ) -> bytes:  
        """Generate speech audio from text."""  
        params = {"text": text}  
        if voice:  
            params["voice"] = voice  
  
        async with httpx.AsyncClient(timeout=120.0) as client:  
            response = await client.get(  
                f"{self.base_url}/tts",  
                params=params  
            )  
            response.raise_for_status()  
        return response.content
```

8.3 Audio Streaming to G1

```
async def send_audio_to_g1(audio_bytes: bytes):  
    """Stream PCM audio to G1's audio_receiver."""  
    url = f"http://172.16.2.242:5050/audio"  
  
    async with httpx.AsyncClient(timeout=30.0) as client:  
        response = await client.post(  
            url,  
            content=audio_bytes,  
            headers={"Content-Type": "audio/wav"}  
        )  
        response.raise_for_status()
```

8.4 Response Chunking

```
async def process_llm_response(response_text: str):
    """Chunk text for low-latency TTS playback."""

    # Split into speakable chunks
    chunks = chunk_text(
        response_text,
        first_chunk_size=50,    # Small first chunk
        rest_chunk_size=150     # Larger rest
    )

    for chunk in chunks:
        audio = await chatterbox_client.generate(chunk)
        await send_audio_to_g1(audio)
```

CHAPTER 9: SLAM NAVIGATION SYSTEM

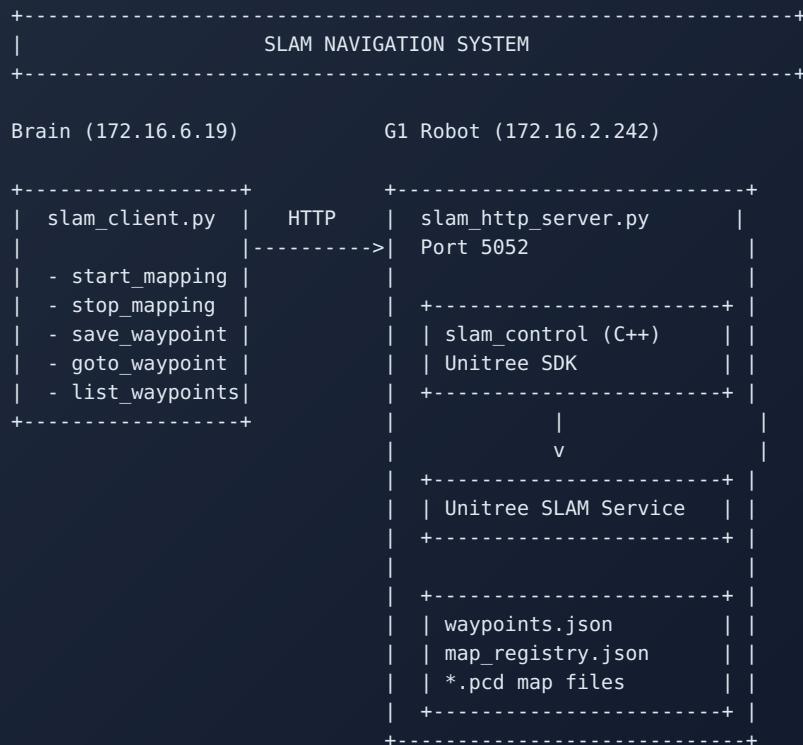
9.1 Overview

The SLAM system enables autonomous mapping and navigation using the Unitree SDK.

Files:

- `slam_client.py` - HTTP client for brain
- `slam_http_server.py` - G1 HTTP server
- `slam_control.cpp` - Unitree SDK binary

9.2 SLAM Architecture



9.3 SLAM API IDs (Unitree SDK)

API ID	Function
1801	START_MAPPING
1802	END_MAPPING
1804	START_RELOCATION
1102	POSE_NAV
1201	PAUSE_NAV
1202	RESUME_NAV

9.4 SLAM Client Methods

```
class SLAMClient:
    def __init__(self, host="172.16.2.242", port=5052):
        self.base_url = f"http://{host}:{port}"

    # Mapping Operations
    def start_mapping(self) -> SLAMResponse
    def stop_mapping(self, map_name: str = None) -> SLAMResponse

    # Navigation Operations
    def start_navigation(self, map_name: str = None) -> SLAMResponse
    def goto_waypoint(self, name: str) -> SLAMResponse
    def pause_navigation(self) -> SLAMResponse
    def resume_navigation(self) -> SLAMResponse

    # Waypoint Management
    def save_waypoint(self, name: str) -> SLAMResponse
    def list_waypoints(self) -> SLAMResponse
    def delete_waypoint(self, name: str) -> SLAMResponse

    # Map Management
    def list_maps(self) -> SLAMResponse
    def delete_map(self, name: str) -> SLAMResponse
```

9.5 Voice Commands for SLAM

Command	Action	Description
"start mapping"	START_MAPPING	Begin SLAM mapping
"stop mapping as office"	STOP_MAPPING:office	Save map as "office"
"start navigation office"	START_NAV:office	Load "office" map
"save this as kitchen"	SAVE WAYPOINT:kitchen	Save current position
"go to kitchen"	GOTO WAYPOINT:kitchen	Navigate to waypoint
"list waypoints"	LIST WAYPOINTS	Show saved waypoints
"pause"	PAUSE_NAV	Pause navigation
"resume"	RESUME_NAV	Resume navigation

CHAPTER 10: G1 ROBOT INTEGRATION

10.1 HTTP Action Bridge

Endpoint: POST `http://172.16.2.242:5051/action`

```
@app.post("/action")
async def execute_action(request: ActionRequest):
    action_name = request.action.upper()

    # Validate action
    if action_name not in VALID_ACTIONS:
        raise HTTPException(400, f"Unknown: {action_name}")

    # Check FSM state
    if not is_fsm_compatible(action_name, get_current_fsm()):
        return {"success": False, "error": "Invalid FSM"}

    # Execute
    if action_name in GESTURE_ACTIONS:
        execute_gesture(action_name)
    elif action_name in MOTION_ACTIONS:
        execute_motion(action_name)

    return {"success": True, "action": action_name}
```

10.2 Unitree API Reference

API ID	Purpose	Example
7101	Set FSM State	SetFsmlId(801)
7104	Height Control	HighStand(), LowStand()
7105	Move Control	Move(vx, vy, vyaw)
7106	Built-in Arm Actions	WaveHand(), ShakeHand()
7108	Custom Arm Actions	ExecuteCustomAction("hug")

API 7106 - Built-in Tasks

ID	Action
0	WaveHand (no turn)
2	ShakeHand (Stage 0)
18	HIGH_FIVE
19	HUG
26	WAVE

API 7108 - Custom Actions

```
namaste1, namaste2, explain1-5, self, self_hug,
battery, Handshake, release_arm
```

10.3 Motion Parameters

```
MOTION_PARAMS = {
    "FORWARD": {"vx": 0.3, "vy": 0, "vyaw": 0, "duration": 2.0},
    "BACKWARD": {"vx": -0.3, "vy": 0, "vyaw": 0, "duration": 2.0},
    "LEFT": {"vx": 0, "vy": 0, "vyaw": 0.3, "duration": 1.5},
    "RIGHT": {"vx": 0, "vy": 0, "vyaw": -0.3, "duration": 1.5},
}
```

10.4 FSM State Transitions

```
ZERO_TORQUE (0)
  |
  | DAMP
  v
DAMP (1)
  |
  | STANDUP
  v
STANDUP (4)
  |
  | READY
  v
READY (801) <-- Arms Enabled
```

Boot Sequence: 0 → 1 → 4 → 801

10.5 Docker Containers on G1

Container	Purpose	Port
audio_receiver	HTTP audio endpoint	5050
tts_audio_player	Speaker playback	-
g1_orchestrator	Main FSM controller	-
arm_controller	Real-time arm (500Hz)	-
http_action_bridge	HTTP to ROS2	5051
talking_gestures	Auto gestures	-
status_announcer	Voice feedback	-
shakal	Face recognition	-

CHAPTER 11: CONFIGURATION SYSTEM

11.1 Main Configuration (config.json)

```
{  
    "g1_audio": {  
        "enabled": true,  
        "host": "172.16.2.242",  
        "port": 5050,  
        "gain": 0.5  
    },  
    "stt_backend": "whisper_server",  
    "tts_backend": "chatterbox",  
    "ollama_model": "naamika:v1",  
    "enable_rag": true,  
    "vad_threshold": 0.5,  
    "silence_duration": 0.5,  
  
    "stt": {  
        "whisper_server": {  
            "host": "172.16.4.250",  
            "port": 8001,  
            "model": "medium"  
        }  
    },  
    "llm": {  
        "deployment_mode": "remote",  
        "remote": {  
            "host": "172.16.4.226",  
            "port": 11434  
        }  
    },  
    "slam": {  
        "host": "172.16.2.242",  
        "port": 5052  
    }  
}
```

CHAPTER 12: API REFERENCE

12.1 Brain Server (172.16.6.19:8080)

Endpoint	Method	Description
/	GET	Web interface
/stream	WebSocket	Audio streaming
/api/health	GET	Health check

12.2 G1 Action Bridge (172.16.2.242:5051)

Endpoint	Method	Body
/action	POST	{"action": "WAVE"}
/health	GET	-

12.3 G1 SLAM Server (172.16.2.242:5052)

Endpoint	Method	Body
/slam/start_mapping	POST	-
/slam/stop_mapping	POST	{"map_name": "office"}
/slam/relocate	POST	{"map_name": "office"}
/slam/goto_waypoint	POST	{"name": "kitchen"}
/waypoint/save	POST	{"name": "kitchen"}
/waypoint/list	GET	-
/map/list	GET	-

12.4 External Services

Service	Endpoint
Chatterbox TTS	GET /tts?text=Hello
Whisper STT	POST /transcribe (multipart)
Ollama LLM	POST /api/chat

CHAPTER 13: DEPLOYMENT GUIDE

13.1 System Requirements

Machine	OS	GPU	Purpose
This PC	Ubuntu 22.04+	Optional	Brain, TTS
Isaac PC	Ubuntu 22.04+	8GB+ VRAM	LLM
New PC	Ubuntu 22.04+	5GB+ VRAM	STT
G1 Robot	ROS2 Humble	-	Control

13.2 Startup Sequence

Step 1: External Services

```
~/scripts/g1_services.sh start chatterbox  
~/scripts/g1_services.sh start ollama  
~/scripts/g1_services.sh start whisper
```

Step 2: G1 Robot

```
ssh unitree@172.16.2.242  
sudo systemctl start slam_server  
cd ~/deployed/v2_1_7 && docker-compose up -d
```

Step 3: Brain

```
cd ~/Downloads/naamika_brain_v2_1_7  
python3 server.py
```

Step 4: Access Open: <https://172.16.6.19:8080>

13.3 SSH Reference

Machine	Command	Password
G1	ssh unitree@172.16.2.242	123
Isaac	ssh isaac@172.16.4.226	7410
New PC	ssh b@172.16.4.250	1997

CHAPTER 14: TROUBLESHOOTING

14.1 Brain Server

Port in use:

```
lsof -i :8080 && kill -9 <PID>
```

14.2 STT Issues

Server not responding:

```
ssh b@172.16.4.250
curl http://localhost:8001/health
```

14.3 LLM Issues

Ollama not responding:

```
ssh isaac@172.16.4.226
OLLAMA_HOST=0.0.0.0:11434 ollama serve
```

14.4 G1 Robot Issues

Docker containers:

```
docker-compose logs -f
docker-compose down && docker-compose up -d
```

FSM stuck:

```
curl -X POST http://172.16.2.242:5051/action \
-d '{"action": "DAMP"}'
```

APPENDIX A: FILE REFERENCE

A.1 Brain Directory

```
naamika_brain_v2_1_7/
├── server.py          # Main server
├── config.json        # Configuration
├── intent_reasoner.py # Intent detection
├── action_registry.py # Actions
├── stop_fast_path.py  # Emergency STOP
├── slam_client.py     # SLAM client
├── naamika_rag.py     # RAG agent
├── chatterbox_tts_client.py
├── vad_pipeline/
│   ├── vad_processor.py
│   ├── audio_recorder.py
│   └── pipeline_manager.py
└── naamika_vectorstore/ # FAISS index
```

A.2 G1 Directory

```
/home/unitree/deployed/v2_1_7/
├── docker-compose.yml
├── config/
└── src/
    ├── g1_orchestrator/
    ├── arm_controller/
    ├── http_action_bridge/
    └── audio_player/

/home/unitree/slam_control/
├── slam_control.cpp
└── slam_http_server.py
```

APPENDIX B: CONFIGURATION VALUES

Parameter	Value
Brain Port	8080
Audio Receiver	5050
Action Bridge	5051
SLAM Server	5052
Chatterbox	8000
Whisper	8001
Ollama	11434
VAD Threshold	0.5
Motion Duration	2.0s
Turn Duration	1.5s

End of Document

NAAMIIKA Voice Assistant System - Technical Documentation v2.1.7

SS Innovations International

Generated: January 27, 2026